

PRODUCT TEARDOWN

PolicyBazaar is India's largest insurance broker, a platform connecting **13.2 Cr users** with **53 insurers** across 800+ products. It earns commission on every policy sold. FY26: ₹6,794 Cr revenue. ₹670 Cr profit. 37% YoY growth.

PolicyBazaar solved distribution. It hasn't solved decision-making.

90%

online insurance
market share

yet only

~1%

of total insurance
is sold online

India's insurance penetration: 3.7% that is half the global average of 7.3%

KEY TAKEAWAY PolicyBazaar dominates online insurance (90%+ share) but digital is stuck at ~1% of total market. The platform shows 80 plans sorted by price but never helps users decide what's right for them. 48% end up under-insured.

PREMIUM FACILITATED

₹29,934 Cr

+42% YoY

FY26 · across health, term, motor, travel

CONVERSION

13.2 Cr

 registered

6.29 Cr bought — only 48% convert

THE BLIND SPOT

48% of buyers pick ₹5L cover

Experts recommend ₹15-25L for metro families.
The platform never tells them the difference.

Sources: PB Fintech FY26 results · IRDAI FY25 · PB 'How India Buys Insurance 2.0'

QUANTITATIVE VALIDATION

What 1,500 negative reviews reveal

2,100 reviews scraped via Python (google-play-scraper). 1,500 negative reviews (1-3 star) categorised by pain point.



Why this matters: 46.7% cite spam calls — the *symptom* of failed digital conversion. PolicyBazaar calls aggressively because the comparison page doesn't convert users on its own. The calls compensate for a product that generates leads, not decisions.

FROM POLICYBAZAAR'S OWN RESEARCH

51%

underestimate illness costs

87%

miscalculate term needs

75%

coverage below ₹10L

QUALITATIVE VALIDATION

"Once you enter data to compare rates, immediately you start receiving calls. The call centre personnel speaks in a dominant way stating you need to complete the process ASAP."

★★★★☆ · 3,101 upvotes · April 2024

"What they tell you before payment is not what happens afterwards. The text on screen will say one thing. The advisor during the call will say another."

★★★★☆ · 467 upvotes · May 2024


What the quotes reveal: Users acknowledge the comparison feature works ("comparing features, price are all good") but the experience after that point breaks down. The platform generates interest but fails to convert it into confident decisions.

ACADEMIC VALIDATION

Iyengar & Lepper (2000): Reducing options from 24 to 6 increased purchase rates from 3% to 30% — a 10x improvement. **PolicyBazaar shows 80.**

KEY TAKEAWAY Five independent sources validate the same pattern : user reviews (1,500 scraped), PolicyBazaar's own report, academic research, competitor benchmarks, and firsthand product walkthrough. The problem isn't opinion. It's structural.

USER JOURNEY: ONBOARDING FIRST TIME USER



Policy Bazar never asks about: income, lifestyle, employer cover, budget, coverage goals and first time buyers

KEY TAKEAWAY: Policy Bazar wins on distribution and scale but no contemporary offers need-based guidance at digital scale. Ditto provides advisory model, while Policy Bazar proves aggregator scale (29,934 Cr.). Nobody combines both which is the structural gap.

COMPETITIVE LANDSCAPE

Where PolicyBazaar wins and what it can learn

Ditto Insurance

Provides: Needs-based advisory, 30-min expert call, 4.9/5 rating

Lacks: Scale (1,055 staff for 800K users), health + term only

PB can learn: Advisory-first model, education before selling, zero spam

InsuranceDekho

Provides: 150K agent network, B2B2C reach into Tier 2/3 cities

Lacks: Same comparison grid, inconsistent agent quality

PB can learn: Agent-assisted model reaches users self-serve can't

Acko

Provides: Simple direct purchase, zero comparison paralysis, own underwriting

Lacks: Limited to own products, not an aggregator, narrow range

PB can learn: Simplicity of zero-choice-overload, clean digital UX

Coverfox

Provides: Same aggregator model, ₹62 Cr revenue

Lacks: No differentiation from PB, much smaller scale

PB can learn: Nothing meaningful — same model, weaker execution

PB platform collects enough to calculate premiums ; not enough to recommend the right coverage.

Problem Framing Canvas

Structuring the problem before solving it

1. What is the problem?

Policy Bazar asks **5 demographic questions** (age, gender, city, members, medical history) and returns **80 different policies sorted by price**. It never asks about income, employer coverage, budget, or life situation, the inputs needed to recommend coverage, not just list it.

Failure A, Decision paralysis: ~52% of registered users never buy (13.2 Cr registered vs 6.29 Cr policies sold, claimed by PB website).

Failure B, Under-insurance: 48% of buyers pick ₹5L or less. Experts recommend ₹15-25L (Source: PB 'How India Buys Insurance 2.0').

2. Who are the customers facing the problem?

First-time health insurance buyers, 25-40, urban India. 65% of PB visitors are under 40. 45% unsure if employer cover is sufficient (PB report). Comfortable buying online but lack financial literacy to evaluate 80 options independently (NBER WP21160).

The abandoner : sees 80 plans, can't choose, leaves, gets spam-called (~52% of users).

The under-buyer : defaults to cheapest ₹5L when they need ₹15-20L (48% of buyers).

3. How do we know it's a problem?

Five independent evidence sources:

① **Firsthand walkthrough :** 7 screens documented. PB asks 5 inputs, shows 80 plans, zero needs assessment.

② **Play Store analysis :** 1,500 negative reviews scraped via Python. 46.7% cite spam calls (symptom of failed digital conversion).

③ **PolicyBazaar's own report :** 48% pick ₹5L or less. 75% below ₹10L. 51% underestimate illness costs. 87% miscalculate term needs.

④ **Competitor validation :** Ditto advisory model: 4.9/5 rating, 800K users. Proves guidance works but doesn't scale.

⑤ **Academic research :** Iyengar & Lepper (2000): 24→6 options increased purchase from 3%→30%. Chernev meta-analysis (99 studies). RAND, NBER confirm insurance-specific choice overload, leading to churn rate.

4. What value is generated by solving this?

For users:

- Adequate coverage (₹12-15L+) instead of cheapest default (₹5L)
- Can decide without spam calls. Confident digital decision removes need for follow-up
- Understand what policy covers before buying because AI explains in plain language

For the business (year-1, assumptions stated):

- New conversions: 10% adoption × 5pp lift × ₹12K premium × 15-20% commission = **₹70-120 Cr**
- Premium uplift: ₹5L → ₹12L avg cover = 2x premium, proportional commission = **₹150-200 Cr**
- Call centre savings: fewer agents if digital converts = **₹30-50 Cr**

Total year-1 impact: ₹250-370 Cr

5. Why should we solve this now?

1. The problem scales with the market : PB premium grew 42% YoY to ₹29,934 Cr (FY26). More users entering the same broken funnel = more wrong decisions at scale.

2. AI makes advisory scalable for the first time: Ditto proved advisory works but needs 1,055 human advisors for 800K users. LLMs deliver conversational assessment at near-zero marginal cost. Not feasible 3 years ago.

3. PolicyBazaar sees the gap but hasn't solved it : Their 'AI Expert' recommends by popularity ('85% of families like yours buy X'), not individual needs. PB Buddy chatbot sits in a corner. They recognise the problem but haven't quite cracked the solution.

4. Regulatory tailwind : IRDAI flagged mis-selling as a concern (FY25 report). Bima Sugam launching 2026. Needs-based advisor aligns with regulatory push toward informed decision-making.

5. Competitive window: No player occupies 'needs-based guidance + self-serve digital'. Ditto = guidance + human. PolicyBazaar = price + self-serve. The empty quadrant is the opportunity.

KEY TAKEAWAY: The problem is structural (not a UX bug), validated by five independent sources, quantifiable (₹250-370 Cr year-1), and timely (AI makes it solvable at scale). It affects both users who abandon and users who buy wrong. Let's hypothesise to navigate to find the root cause.

HYPOTHESIS FORMATION AND SELECTION

From validated problem to testable hypothesis

CHOSEN

H1 No needs assessment causes decision paralysis and under-insurance

PB asks 5 demographic questions, shows 80 plans, never assesses income, employer cover, or budget. Result: **52% abandon** (13.2Cr registered vs 6.29Cr bought).

48% of buyers pick ₹5L or less against ₹15-25L recommended (PB report). This is the **root cause** the other two hypotheses are downstream effects of this one.

H2 Spam calls destroy trust which is a symptom of failed digital conversion

46.7% of negative reviews cite unwanted calls (Play Store scrape, 701 of 1,500 reviews). But PB calls *because* the comparison page doesn't convert users on its own.

This is a **symptom** of the primary hypothesis, not an independent problem. Fixing **H1** partially fixes this.

H3 Misleading info and post-sale abandonment erode trust

8.9% cite misleading information; 11% cite poor claim support (Play Store scrape).

"What they tell you before payment is not what happens afterwards" (**2★, 467 upvotes**).

This is **downstream** of H2; pushy agents over-promise to close sales.

RICE PRIORTISATION- Why each hypothesis scored what it did

Factor	H1: No needs assessment	H2: Spam calls	H3: Trust erosion
Reach	8 13.2Cr users. 52% don't convert + 75% of buyers under-insured.	9 46.7% of negative reviews = largest single complaint. Anyone who enters phone number.	4 8.9% cite misleading, 11% poor claims. Real but smaller segment.
Impact	9 Paradigm shift: compare-by-price → recommend-by-fit. Iyengar: 24→6 options = 3%→30% purchase.	3 Removes annoyance but doesn't help users decide. 80 plans without calls is still 80 confusing plans.	4 Better transparency helps but users still pick wrong policy, root cause unchanged.
Confidence	9 5 independent sources converge: walkthrough, 1,500 reviews, PB report, Ditto validation, academic research.	7 Symptom confirmed (46.7%). But don't know conversion impact if calls stopped. PB may have data showing calls drive revenue.	6 Quotes exist but unclear if agents or platform cause the misleading info. If agents, product fix won't help.
Effort	3 Buildable: React + Groq API + Vercel. Already built Claude Reflect, SBI MF chatbot. Core work = system prompt.	9 Can't build a solution. Don't control PB's call centre or CRM. A 'do not call' toggle isn't a product.	4 Policy explainer is buildable but thin as portfolio piece. Less impressive than conversational AI.
Score	216	21	24

KEY TAKEAWAY H1 scores 10× higher than alternatives. It's the root cause (fixing it partially fixes H2 and H4), has the highest confidence (5 independent sources), and is buildable as a working prototype. Selected for solution design.

SOLUTION EVALUATION AND SELECTION

Four approaches to solve H1- one clear winner

S1 AI Conversational Advisor



Chat-based interface asks about your life (income, family, employer cover, health, budget) and recommends specific coverage with 2-3 best-fit policies and plain-language reasoning.

Why it wins: Closest to Ditto's proven advisory model (4.9/5) but delivered digitally at scale. Solves both H1 (decision paralysis) and H3 (under-insurance). Occupies the empty quadrant — needs-based guidance + self-serve digital — where no competitor sits.

S2 Smart Coverage Calculator

Structured form : user inputs income, city, family size, employer cover. Tool calculates recommended coverage and shows 3 matching policies.

Solves: H3 (under-insurance): gives the right coverage number.

Misses: H1 partially, as user still faces a form, not a conversation. Can't handle nuance ('I have diabetes but it's managed'). Feels like another PolicyBazaar form, not a paradigm shift.

S3 AI Policy Simplifier

User sets preferences (no room rent limit, low co-pay). AI scores all plans against preferences and surfaces top 3 with explanations.

Solves: H1 partially, as reduces 80 plans to 3.

Misses: Assumes users know what preferences matter. A first-time buyer doesn't know if room rent limit is important. Doesn't assess how much coverage is needed,

S4 Smart Coverage Calculator

Step-by-step visual flow teaching insurance concepts (co-pay, room rent, sub-limits) and narrowing options based on answers. Part education, part decision tool.

Solves: The literacy gap. NBER confirms insurance literacy predicts better choices.

Misses: Takes 10-15 minutes (high drop-off risk). Doesn't give a direct recommendation Content product, not an AI product.

RICE PRIORTISATION- Why each hypothesis scored what it did

Factor	S1: AI Conversational Advisor	S2: Smart Coverage Calculator	S3: AI Policy Simplifier	S4: Smart Coverage Calculator
Reach	8 Same user base. Both abandoners and under-buyers benefit from conversational guidance.	8 Same user base. Calculator serves same segment.	7 Only users who already know preferences. First-timers don't.	5 10-15 min flow. Many drop off before finishing.
Impact	9 Full paradigm shift. Ditto proved it at 4.9/5. Solves both H1 and H3 completely.	5 Gives right amount but not right policy. Form, not a conversation. Partial fix.	4 Reduces options but user still picks cover amount blindly. H3 unsolved.	4 Educates but doesn't recommend. User still decides alone at the end.
Confidence	9 Ditto validated advisory model in-market. Iyengar confirms fewer options = better outcomes.	6 No market validation of calculator-only approach. Calculators exist but adoption unclear.	5 No validation. Assumes users can articulate what matters. They can't (NBER).	6 NBER confirms literacy helps. But no evidence this specific format drives purchases.
Effort	3 React + Groq API + Vercel. Core work and system prompt engineering. Already built similar.	2 Simpler. No LLM needed, deterministic logic. Faster to build and test.	3 Needs LLM for scoring. Similar effort to S1 but thinner outcome.	2 Mostly content. Minimal AI. Simpler build but less impressive.
Score	216	120	47	60

KEY TAKEAWAY H1 scores 10× higher than alternatives. It's the root cause (fixing it partially fixes H2 and H4), has the highest confidence (5 independent sources), and is buildable as a working prototype. Selected for solution design.

AI Insurance Advisor- how it works

THE SOLUTION IN ACTION

BEFORE (PolicyBazaar today): 5 demographic questions → 80 plans sorted by price → zero guidance → user abandons or picks cheapest

AFTER (AI Advisor): 7 life-situation questions → specific coverage amount with reasoning → 2-3 best-fit policies → user decides with confidence

How each step maps to the problem it solves

1. User selects insurance type

Health or Term Life. Sets context for the entire conversation flow.

Solves: PB jumps straight to forms. No framing of what you're looking for.

2. AI asks about your life, not just demographics

Income, family structure, city, employer cover, health conditions, budget. 7 questions asked conversationally, one at a time.

Solves: PB asks 5 demographic inputs. Never asks income, employer cover, or budget.

3. AI calculates your specific coverage need

Uses insurance heuristics (health: 3-5× income adjusted for city + family. Term: 10-15× income minus existing cover + loans).

Solves: PB shows a dropdown, user picks a number with no basis. 48% pick ₹5L.

4. Shows 2-3 best-fit policies with reasoning

Each policy card explains WHY it fits this user. 'This plan has no room rent limit, which matters because you're in Delhi where hospital costs are high.'

Solves: PB shows 80 undifferentiated plans. No 'why this one for you.'


5. User gets complete recommendation — no phone required

Full coverage amount + policy picks + reasoning delivered in one session. Optional: schedule callback or download summary.

Solves: PB gates plans behind phone capture → spam calls → 46.7% complaint rate.

KEY TAKEAWAY This isn't a mockup; it's a working AI advisor deployed on Vercel. Every step directly addresses a validated problem: the 7 questions PB never asks, the coverage calculation PB never does, and the plain-language reasoning PB never provides. The user walks away with a recommendation, not a homework assignment of 80 plans.

Live Link: [PolicyBazaar AI Advisor — Find the Right Insurance for Your Life](#)



Demo Scenario:

28-year-old, Delhi, ₹12L income, no employer cover, family of 3, no pre-existing conditions

→ **AI recommends ₹15L health cover with 3 policies + reasoning**

How we measure whether the solution actually works

NORTH STAR

Coverage Adequacy Rate

% of users who purchase coverage at or above the AI-recommended amount

Why this metric: It captures the core problem we're solving. Currently, 48% of buyers pick ₹5L or less against a recommended ₹15-25L (PB 'How India Buys Insurance 2.0'). If the advisor works, this gap closes and users buy what they actually need, not what's cheapest.

Current baseline: ~25% adequate (only 25% of buyers currently choose ₹10L+ cover - PB report). **Target:** 60%+ adequate coverage within year 1. Benchmark: Ditto achieves near-100% adequacy because advisors calculate needs. Our AI replicates this digitally.

LEADING INDICATORS : early signals that predict North Star movement

Recommendation Completion Rate

What: % of users who complete all 7 questions and receive a full recommendation.
Why it matters: Currently ~52% of PB registered users never buy (13.2Cr vs 6.29Cr — PB footer). If users drop off mid-conversation, the advisor isn't engaging enough.
Target: 70%+ completion. Baseline: typical chatbot completion is 40-50% (industry avg).

Recommendation Acceptance Rate

What: % of users who click through to a recommended policy after receiving the AI's recommendation.
Why it matters: Measures whether users trust the AI's guidance enough to act on it. If acceptance is low, the reasoning isn't convincing or the policies don't match expectations.
Target: 40%+ acceptance. PB's current quote-to-purchase is ~48% but with spam call assistance.

Average Recommended Coverage

What: Mean coverage amount the AI recommends across all users. Validates that heuristics are working correctly.
Why it matters: Current avg is ~₹5L (48% of buyers - PB report). If AI consistently recommends ₹12-15L+, the heuristics are correcting the under-insurance gap.
Target: ₹12-15L average. Red flag if below ₹8L (heuristics too conservative) or above ₹25L (over-recommending).

Time to Recommendation

What: Median time from first question to complete recommendation delivery.
Why it matters: Ditto takes 30 minutes with a human advisor. Our AI should deliver the same quality in a fraction of the time — that's the scalability advantage.
Target: Under 5 minutes. Red flag if over 8 min (conversation too long, users will abandon).

GUARDRAILS : what we watch to ensure we're not causing harm

Recommendation Accuracy

Do insurance experts agree with AI recommendations? Spot-check 50 recommendations/month against expert assessment.
Threshold: 85%+ agreement. Below this, heuristics need recalibration. Benchmark: Ditto's human advisors as the quality standard.

Premium Affordability

Is the AI recommending coverage users can actually afford? The budget question (step 7 in our flow) constrains this, but we verify: recommended annual premium should be under 5% of stated income. **Threshold: 95%+ of recommendations within budget.**

Edge Case Escalation Rate

Complex cases (NRIs, senior citizens, chronic conditions) should be flagged for human advisor review, not handled by AI alone.
Threshold: 100% of identified edge cases get escalation prompt. IRDAI distinguishes advice (licensed) from comparison (allowed) — AI must stay on the right side.

KEY TAKEAWAY The North Star directly measures the problem we set out to solve — under-insurance. Leading indicators track whether users engage, trust, and act on recommendations. Guardrails ensure the AI doesn't cause harm through inaccurate, unaffordable, or out-of-scope advice. Every metric has a baseline from our research and a defensible target.

What could go wrong and what I learned

Failure Mode	Why It Matters	Mitigation
AI recommends insufficient or excessive coverage	Heuristics use general rules (3-5× income for health, 10-15× for term) not actuarial precision. Edge cases : high-cost cities, rare conditions, may produce incorrect recommendations. User harm if they follow bad advice.	Expert spot-checks (85%+ accuracy guardrail from Slide 8). Range validation: flag any recommendation below ₹5L or above ₹50L for review. Mandatory disclaimer: 'AI-generated guidance, not licensed financial advice.'
Spam calls persist even with better digital conversion	46.7% of negative reviews cite calls (our Play Store scrape). PB may continue calling because it's profitable and we don't control their call centre or sales incentives. Solving H1+H3 reduces the NEED for calls but can't enforce call discipline.	Acknowledged as an organisational change, not a product change. The AI advisor gives users information advantage as they know what they need before anyone calls, making them harder to mis-sell to. This is a mitigation, not a fix.
Regulatory risk: IRDAI 'advice' vs 'comparison'	IRDAI distinguishes between insurance comparison (allowed for brokers) and insurance advice (requires advisory licence). Our AI recommends specific coverage amounts and policies , this could be classified as advice. IRDAI flagged mis-selling in FY25 report.	Frame as 'guidance' not 'advice.' Clear disclaimers on every screen. Escalate complex cases to licensed advisors. Monitor regulatory developments. Bima Sugam (IRDAI marketplace launching 2026) may reshape rules.
Users don't trust AI for high-stakes financial decisions	Insurance involves significant money (₹15-25K/year for health, more for term life). Users may hesitate to trust a chatbot's recommendation for a decision that affects their family's financial safety. Trust deficit is real and 8.9% of PB reviews cite misleading info.	Show reasoning transparently ('I recommend ₹15L because your income is ₹12L and you're in Delhi where hospital costs are high'). Offer 'Schedule a callback' for human confirmation. Cite claim settlement ratios and policy details.
Business model conflict: fewer displayed policies = less commission surface	PB's revenue = f(policies sold × premium × commission). Showing 2-3 instead of 80 reduces opportunities. Innovator's dilemma: PB won't cannibalise their high-listing model even if it's better for users.	Counter-argument: higher coverage = higher premium = higher commission per sale. ₹15L policy earns ~3× more commission than ₹5L. Better conversion rate (fewer abandonments) offsets fewer displayed options. Year-1 SOM:
Conversation goes off-track or user asks unrelated questions	LLMs can be unpredictable. User might ask about stock markets, argue with recommendations, or provide nonsensical inputs. Degrades experience and trust.	Constrained system prompt with strict conversation flow. Graceful redirects ('That's outside what I can help with , let's get back to finding the right coverage for you'). Fallback to structured form if conversation fails.

REFLECTION

What this teardown taught me

On problem validation: The original hypothesis was 'too many options cause confusion.' The data told a different story, only 3.3% explicitly mentioned confusion. But 46.7% cited spam calls, which led to a stronger insight: the calls are the symptom of a platform that generates leads, not decisions. Following the data changed the narrative.

On building vs theorizing: Building the prototype forced design decisions that slide decks skip. How does the AI handle 'I don't know my income'? What if someone has both diabetes and hypertension? These edge cases only surface when you build.

Limitations I acknowledge

- No primary user interviews conducted : validation is from reviews + reports, not direct conversation.
- Heuristics are generalisations, not actuarial models. Edge cases will fail.
- Prototype uses Groq free tier: production would need more robust infrastructure.
- Under-insurance data is from PB's own report: possible selection bias.

What I'd do next

- Conduct 15-20 user interviews to validate the conversational flow.
- A/B test AI advisor vs standard comparison page on a subset of PB traffic.
- Partner with an IRDAI-licensed advisor to validate heuristics.
- Build term life insurance flow (prototype currently handles health + term).
- Explore integration with PolicyBazaar's API for real-time policy data.